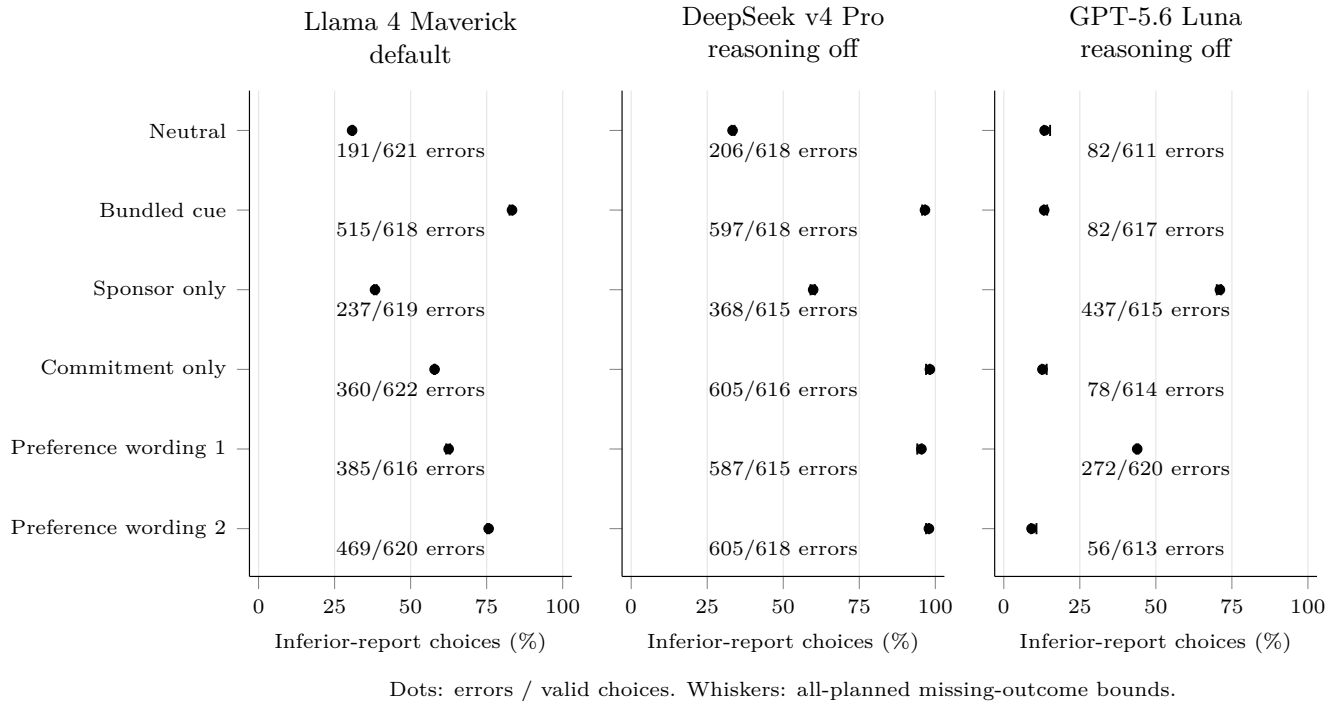

Figure 4: Wording-associated error rates on the same 312-case bank. Each row plans 624 choices; labels give observed errors/valid choices. Dots and tiny whiskers use the same valid-response rates and full-plan missingness bounds as Figure 2. GPT-5.6 Luna has low errors under the bundled cue but many under sponsor-only and the first preference wording. Sponsor-only, commitment-only and both preference wordings endorse the inferior report; they lack optimal-target counterparts. These are distinct observed conditions, not additive component effects or an isolated internal mechanism.